

TRANSFORMER ARCHITECTURE

Vincent Guigue,
inspiré des supports de Benjamin Piwowarski & Thomas Gérald



INTRODUCTION



Aggregation problem

Sequence : $\mathbf{s} = \{s_1, \dots, s_L\}$

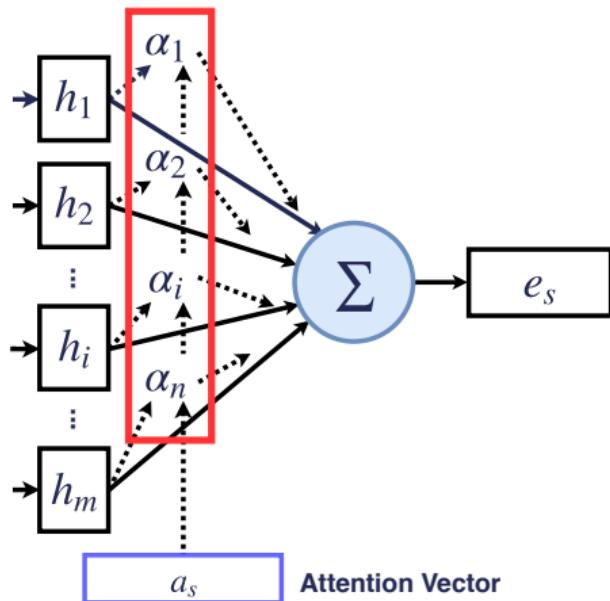
Representation learning : $s_i \rightarrow \mathbf{e}_i \in \mathbb{R}^Z$

Decision on $\mathbf{s} \Rightarrow$ Representation / mapping $R(\mathbf{s}) : \mathbb{R}^{L \times Z} \mapsto \mathbb{R}^d$

+ Decision function $f(R(\mathbf{s})) : \mathbb{R}^d \mapsto \mathcal{Y}$

- RNN + last hidden layer (ext: Bi-RNN)
- CNN/RNN + pooling
- CNN/RNN + attention = learnt weighted sum

Attention & aggregation vs Diffusion process



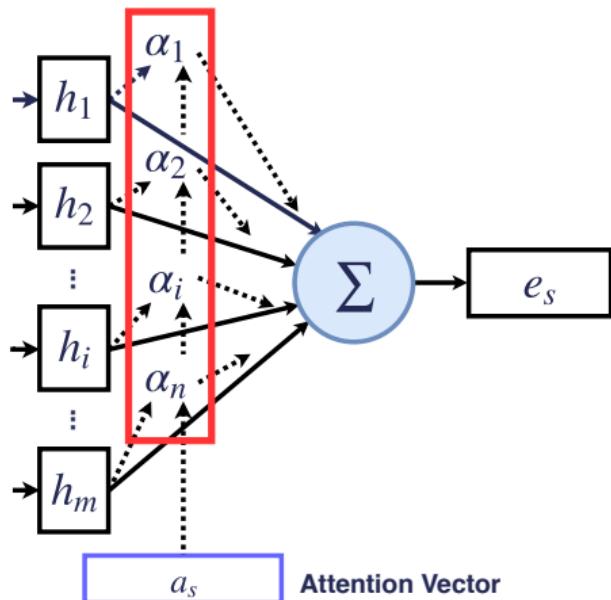
- Learning to weight localized representation
- Many efficient applications [2014-2018]
- Few parameters
- Attention vector = Query vector
- How to code it in pytorch?

Practical Session:

Q1 Preliminary tests: uniform attention

Q2 Classical global attention

Attention & aggregation vs Diffusion process



- Learning to weight localized representation
- Many efficient applications [2014-2018]
- Few parameters
- Attention vector = Query vector
- How to code it in pytorch?

Q3 \Rightarrow Reconstruction of **all embedding** by recombination = diffusion process



Inspiration from memory networks (question answering)

Task:

Sam walks into the kitchen.
Sam picks up an apple.
Sam walks into the bedroom.
Sam drops the apple.

Q: Where is the apple?
A. Bedroom

Brian is a lion.
Julius is a lion.
Julius is white.
Bernhard is green.

Q: What color is Brian?
A. White

Mary journeyed to the den.
Mary went back to the kitchen.
John journeyed to the bedroom.
Mary discarded the milk.

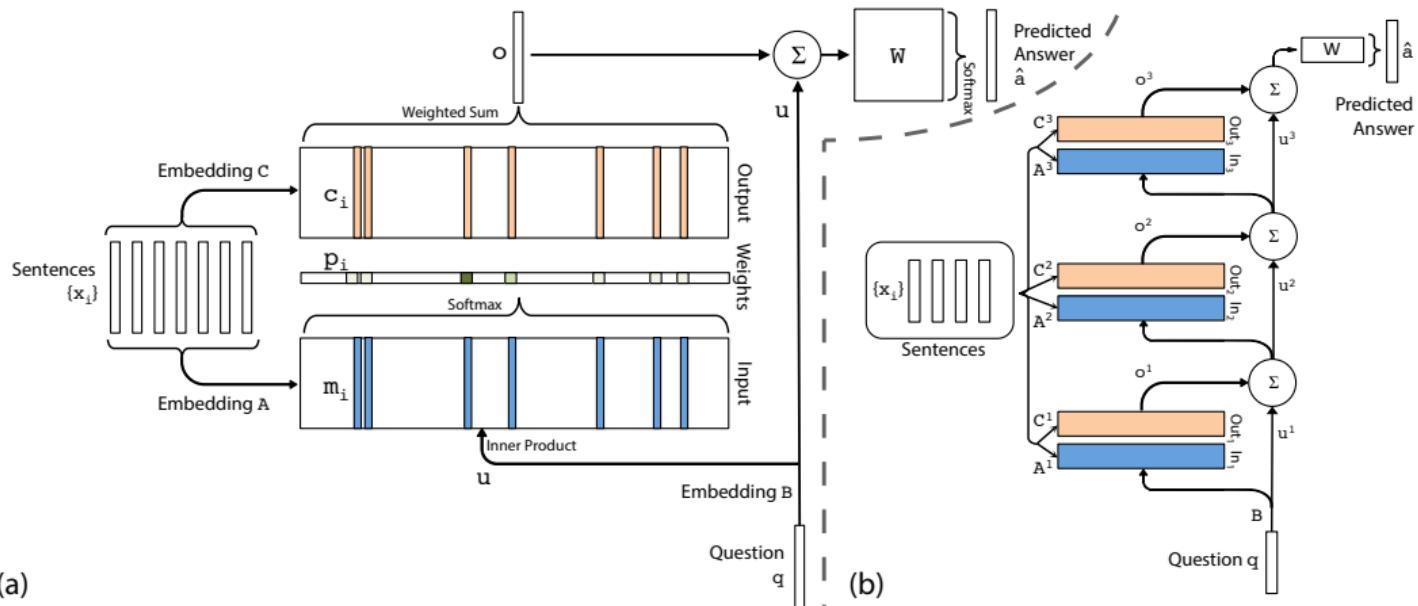
Q: Where was the milk before the den?
A. Hallway

Babi dataset

How to anchor question answering in Knowledge Bases?

Inspiration from memory networks (question answering)

How to anchor question answering in Knowledge Bases?



Inspiration from memory networks (question answering)

How to anchor question answering in Knowledge Bases?

Story (1: 1 supporting fact)	Support	Hop 1	Hop 2	Hop 3
Daniel went to the bathroom.		0.00	0.00	0.03
Mary travelled to the hallway.		0.00	0.00	0.00
John went to the bedroom.		0.37	0.02	0.00
John travelled to the bathroom.	yes	0.60	0.98	0.96
Mary went to the office.		0.01	0.00	0.00
Where is John? Answer: bathroom Prediction: bathroom				

Story (2: 2 supporting facts)	Support	Hop 1	Hop 2	Hop 3
John dropped the milk.		0.06	0.00	0.00
John took the milk there.	yes	0.88	1.00	0.00
Sandra went back to the bathroom.		0.00	0.00	0.00
John moved to the hallway.	yes	0.00	0.00	1.00
Mary went back to the bedroom.		0.00	0.00	0.00
Where is the milk? Answer: hallway Prediction: hallway				

Story (16: basic induction)	Support	Hop 1	Hop 2	Hop 3
Brian is a frog.	yes	0.00	0.98	0.00
Lily is gray.		0.07	0.00	0.00
Brian is yellow.	yes	0.07	0.00	1.00
Julius is green.		0.06	0.00	0.00
Greg is a frog.	yes	0.76	0.02	0.00
What color is Greg? Answer: yellow Prediction: yellow				

Story (18: size reasoning)	Support	Hop 1	Hop 2	Hop 3
The suitcase is bigger than the chest.	yes	0.00	0.88	0.00
The box is bigger than the chocolate.		0.04	0.05	0.10
The chest is bigger than the chocolate.	yes	0.17	0.07	0.90
The chest fits inside the container.		0.00	0.00	0.00
The chest fits inside the box.		0.00	0.00	0.00
Does the suitcase fit in the chocolate? Answer: no Prediction: no				



Conclusion

Current limitations with RNN:

- Long dependencies
- Global message representation
 - Continuous indexing
- Targeted dependencies
 - Co-reference resolution

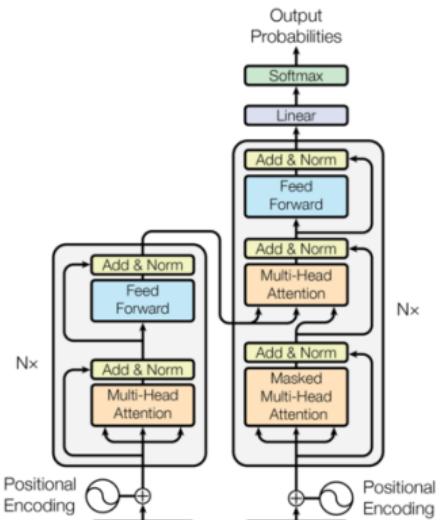
ARCHITECTURE



Transformer generality

Two sources to understand Transformers [easier than the original article]

- J. Alammar 2018
<http://jalammar.github.io/illustrated-transformer/>
- Alexander Rush 2018
<http://nlp.seas.harvard.edu/2018/04/03/attention.html>
- P. Bloem 2019 <http://www.peterbloem.nl/blog/transformers>

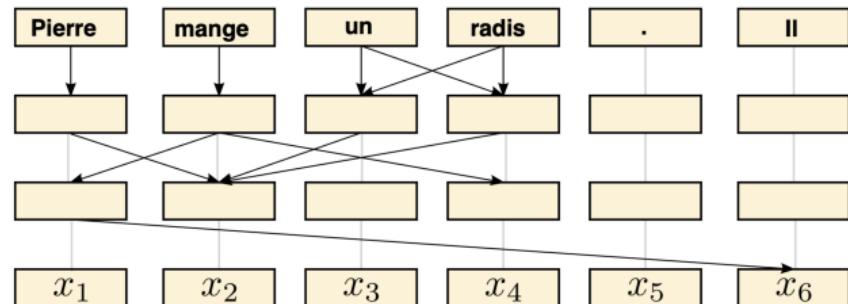


Overcoming classical pictures
⇒ re-implementing a transformer block

Transformer generality

General idea: representation diffusion + learning the diffusion

Scale choice



Character
+ (very)
small
vocabulary
- no
semantics

Word
+ Semantics
++
- Huge
vocabulary
- Unknown
words

[book] Sennrich, R. et al. 2016. Neural Machine Translation of Rare Words with Subword Units.

les chats sont des félidés

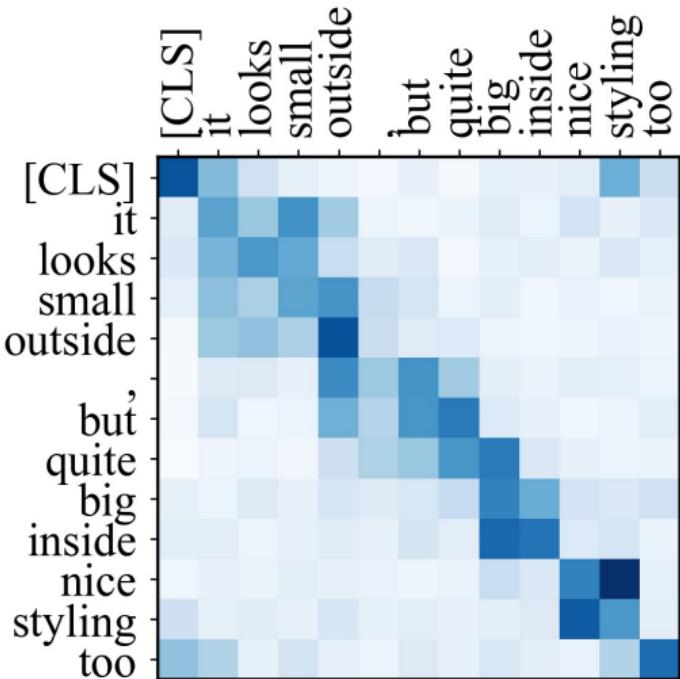
est transformé en

[CLS] _les _chats _sont _des _f éli dés [SEP]

Attention block

- 1 Computing self-attention between tokens
- 2 Normalizing attention (softmax)
- 3 Exploiting attention to recombine token representation

⇒ How to compute interactions
(=self attention)



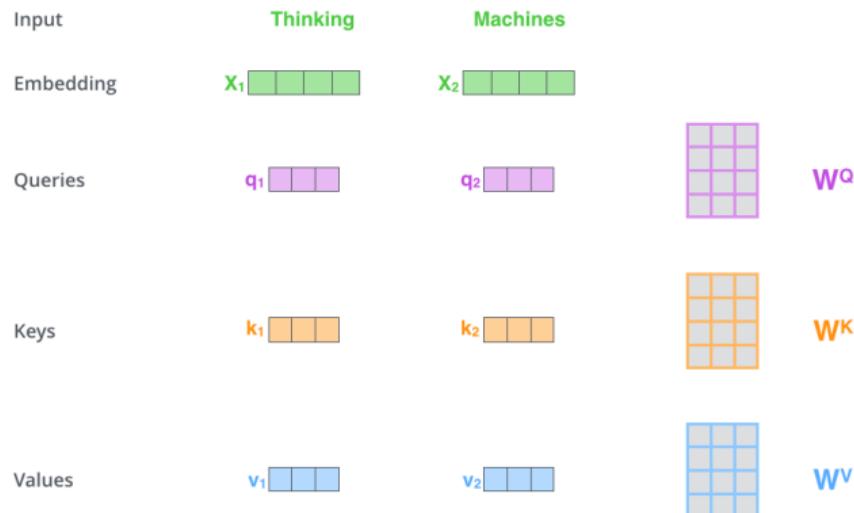


Attention block

Interactions through vectors derived from the embedding

- 1 Compute Keys, Queries, Values by linear mapping

Note: here are the main parameters of the transformer architecture

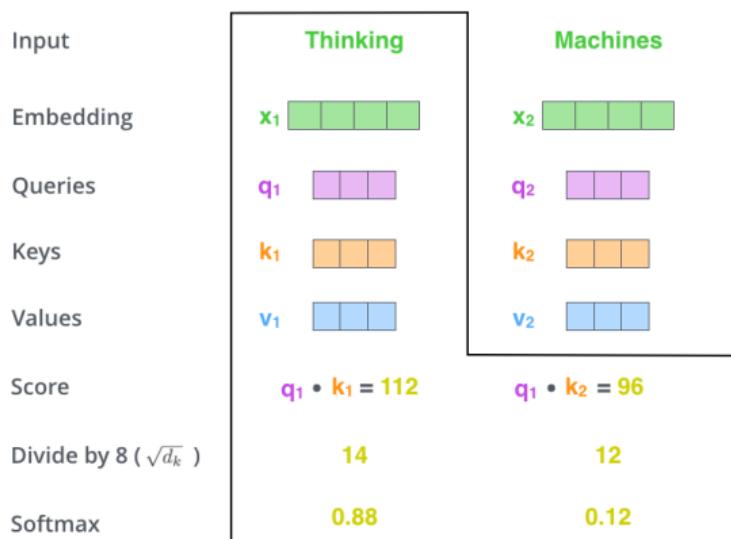




Attention block

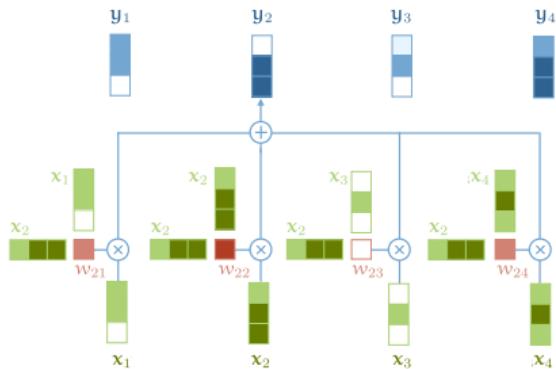
Computing the attention

- 1 Inner product between keys/queries
- 2 Division by key/query size
 - Softmax too sensitive to large values
 - Temperature
- 3 Softmax \Rightarrow over the relevant dimension



Attention block

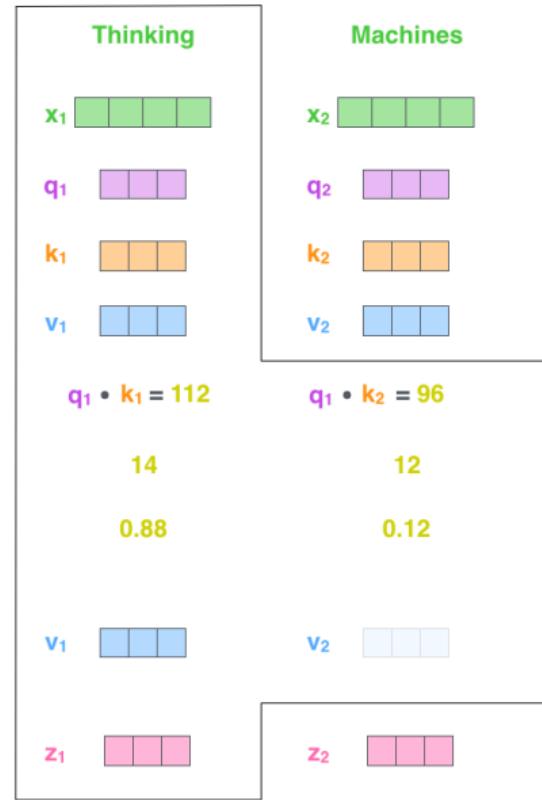
New values:



$$v_i^{\ell+1} = \sum \alpha_k v_k^\ell$$

How to come back to the embedding space?
 $\text{value} \mapsto \text{embedding}$

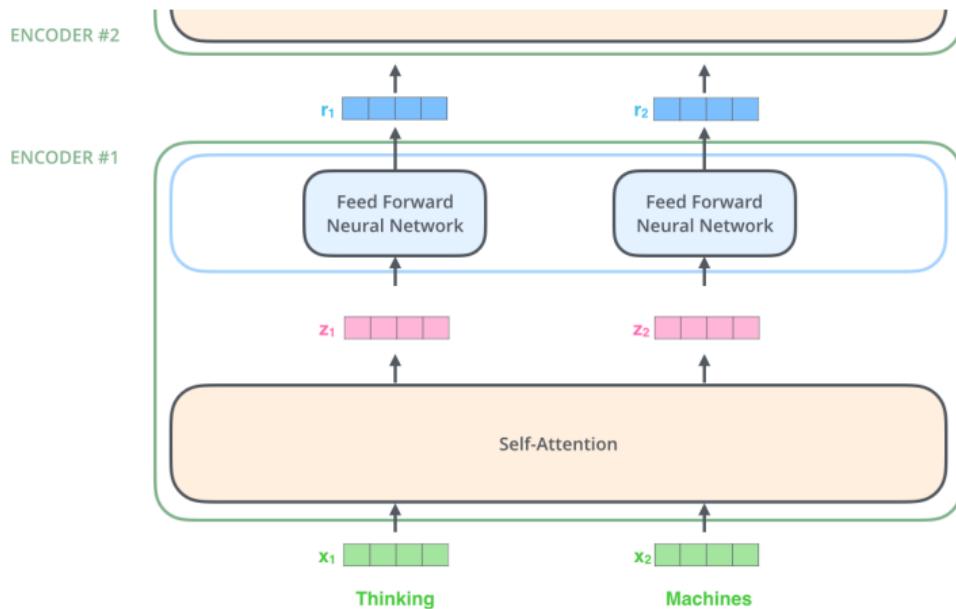
Input
Embedding
Queries
Keys
Values
Score
Divide by 8 ($\sqrt{d_k}$)
Softmax
Softmax X Value
Sum





Attention block

Use a fully connected to map to embedding space



Parallelism

All tokens at once !

$$\mathbf{X} \quad \mathbf{W}^Q \quad \mathbf{Q}$$

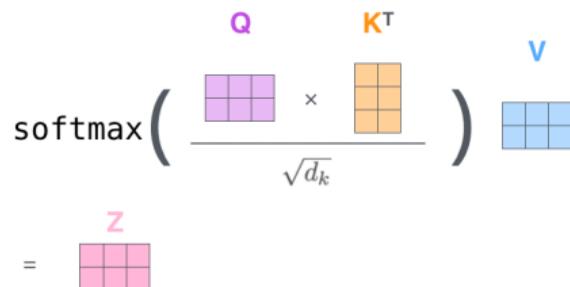

A diagram illustrating the first step of parallel computation. On the left, a green 3x3 matrix labeled \mathbf{X} is multiplied by a purple 3x4 matrix labeled \mathbf{W}^Q . The result is a purple 3x3 matrix labeled \mathbf{Q} .

$$\mathbf{X} \quad \mathbf{W}^K \quad \mathbf{K}$$


A diagram illustrating the second step of parallel computation. On the left, the same green 3x3 matrix \mathbf{X} is multiplied by an orange 3x4 matrix labeled \mathbf{W}^K . The result is an orange 3x3 matrix labeled \mathbf{K} .

$$\mathbf{X} \quad \mathbf{W}^V \quad \mathbf{V}$$


A diagram illustrating the third step of parallel computation. On the left, the same green 3x3 matrix \mathbf{X} is multiplied by a blue 3x4 matrix labeled \mathbf{W}^V . The result is a blue 3x3 matrix labeled \mathbf{V} .

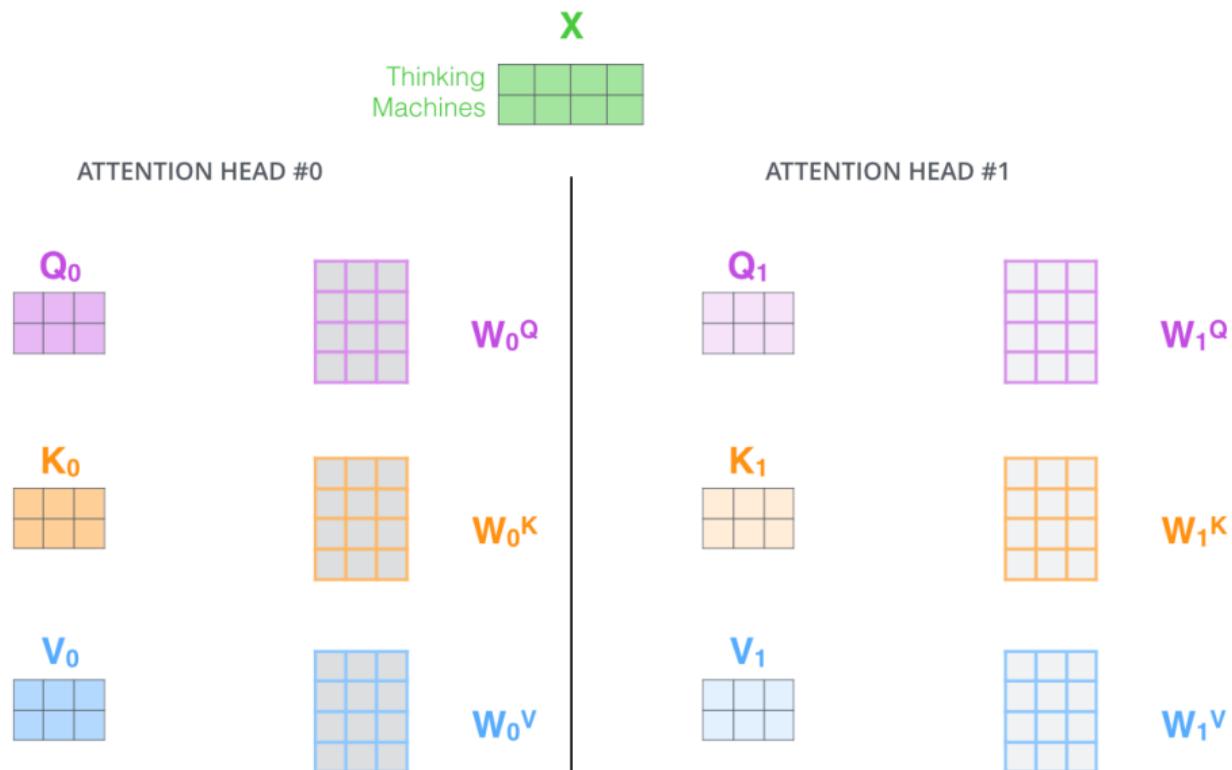
$$\text{softmax}\left(\frac{\mathbf{Q} \times \mathbf{K}^T}{\sqrt{d_k}} \right) \mathbf{V} = \mathbf{z}$$


A diagram illustrating the final step of parallel computation. It shows the softmax function applied to the product of \mathbf{Q} (purple 3x3 matrix) and \mathbf{K}^T (orange 3x3 matrix), divided by $\sqrt{d_k}$. The result is then multiplied by a blue 3x3 matrix labeled \mathbf{V} . The final output is a pink 3x3 matrix labeled \mathbf{z} .



Multiple heads

Why only **one** interaction?



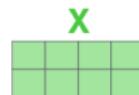


Multiple heads

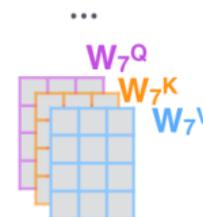
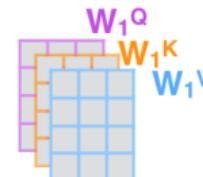
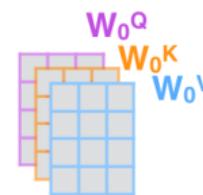
Why only one interaction?

1) This is our input sentence*
each word*

Thinking
Machines



2) We embed each word*
3) Split into 8 heads.
We multiply X or R with weight matrices



4) Calculate attention using the resulting Q/K/V matrices



5) Concatenate the resulting Z matrices, then multiply with weight matrix W^O to produce the output of the layer



W^O



* In all encoders other than #0, we don't need embedding.
We start directly with the output of the encoder right below this one





Multiple heads

Why only **one** interaction?

Typically, values are concatenated then mapped to rebuild embeddings

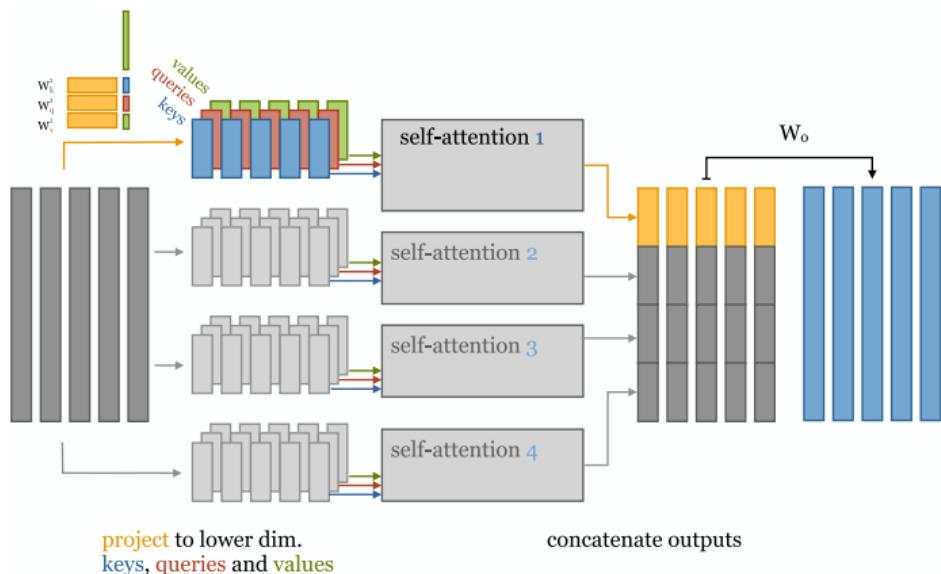
BERT

- Embedding : 768
- key/query/values : 64
- ... & 12 heads

$$\Rightarrow 12 \times 64 = 768$$

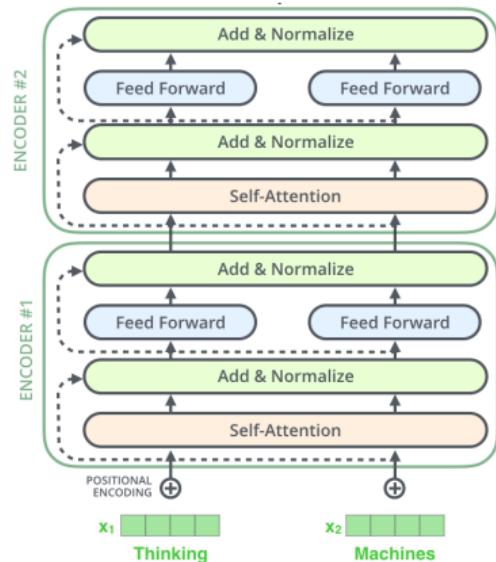
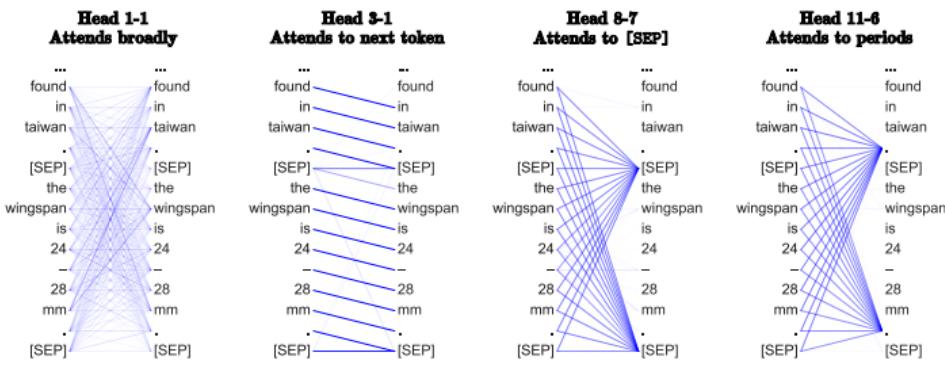
ChatGPT:

- Embedding : 12288
- key/query/values : 128
- 96 heads



A Multiple layers

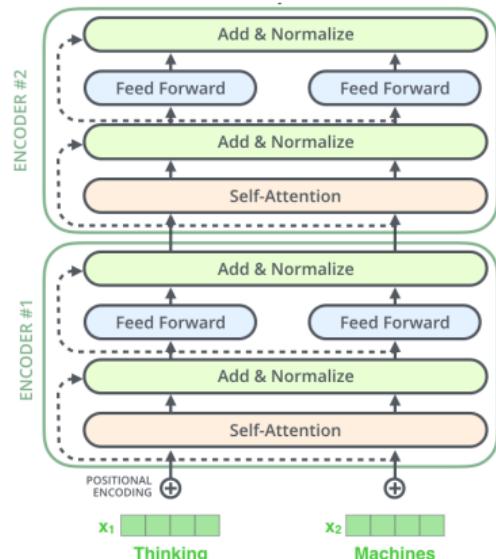
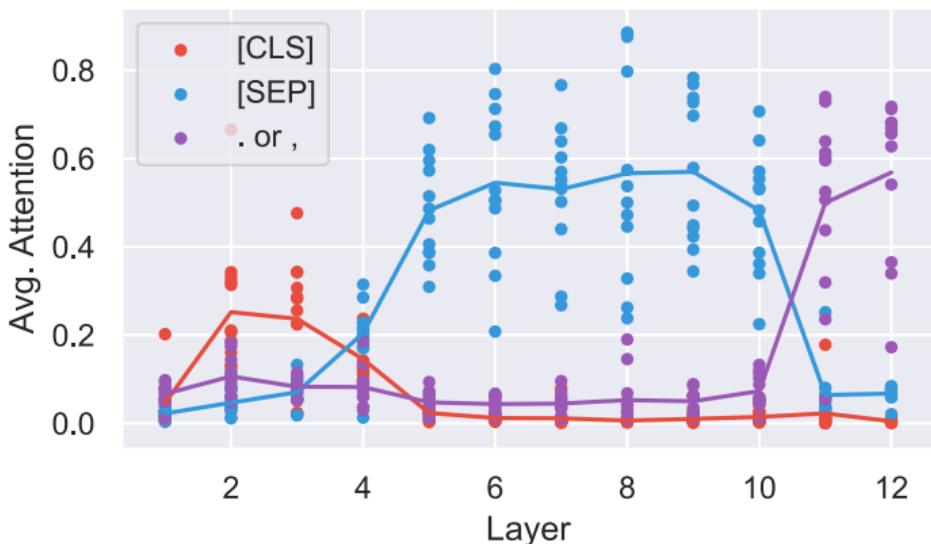
- Diffusion is made step-by-step
- Bertology shows that different relations are extracted at different scales





Multiple layers

- Diffusion is made step-by-step
- Bertology shows that different relations are extracted at different scales

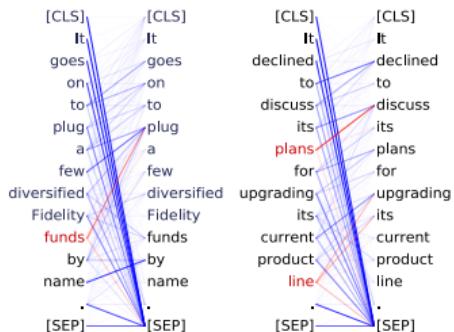


Multiple layers

- Diffusion is made step-by-step
- Bertology shows that different relations are extracted at different scales

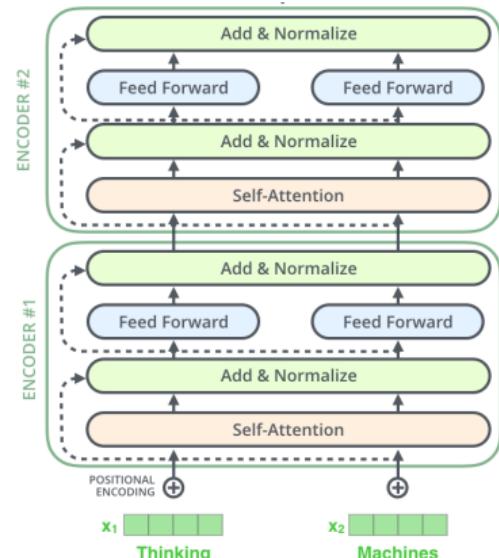
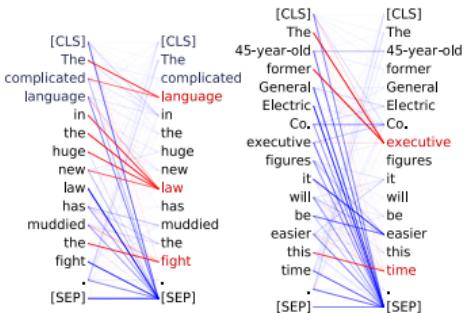
Head 8-10

- Direct objects attend to their verbs
- 86.8% accuracy at the `dobj` relation



Head 8-11

- Noun modifiers (e.g., determiners) attend to their noun
- 94.3% accuracy at the `det` relation





Position coding

Until now, Transformer does not take position into account !!

POSITIONAL
ENCODING

0	0	1	1
---	---	---	---

0.84	0.0001	0.54	1
------	--------	------	---

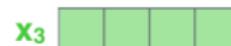
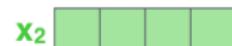
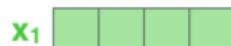
0.91	0.0002	-0.42	1
------	--------	-------	---

+

+

+

EMBEDDINGS



INPUT

Je

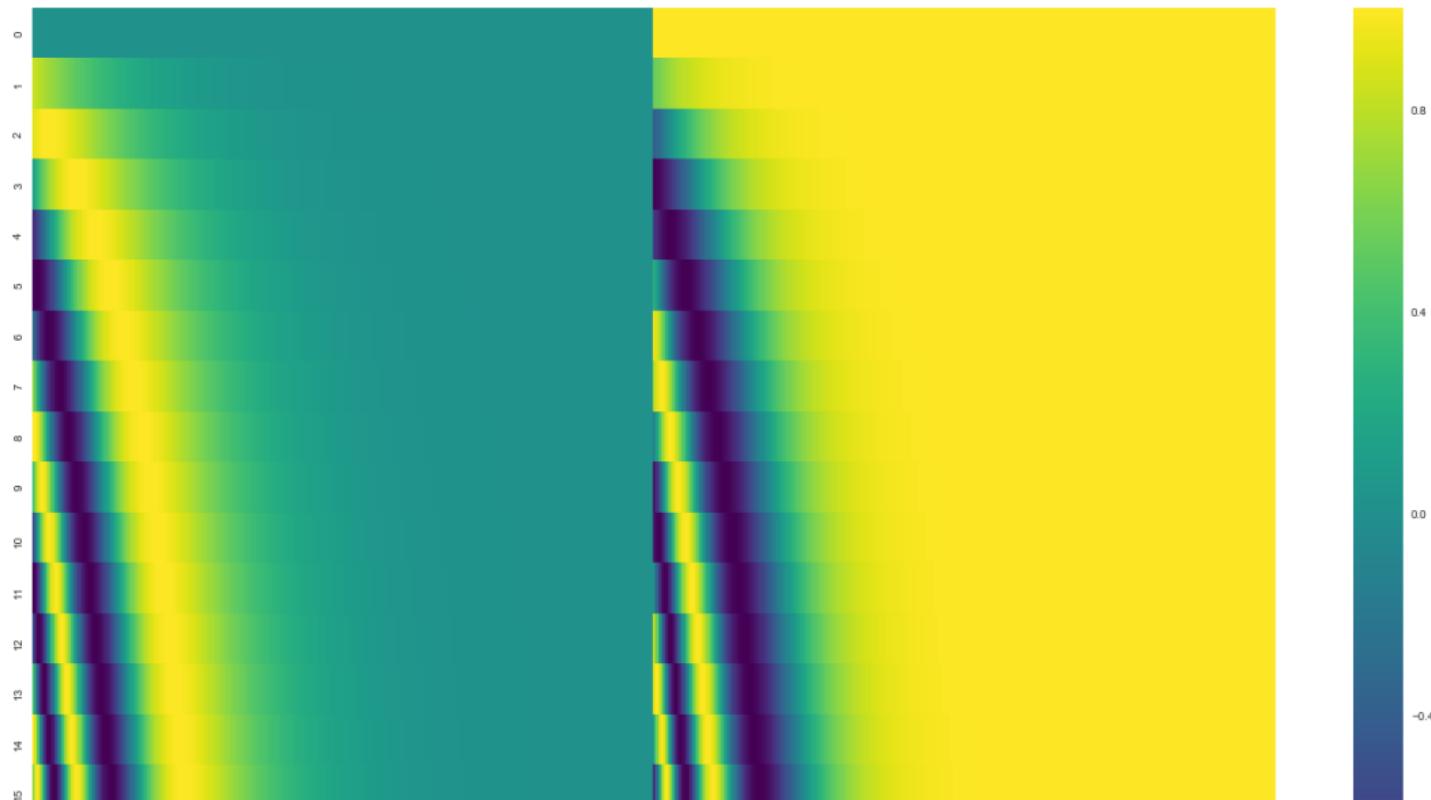
suis

étudiant



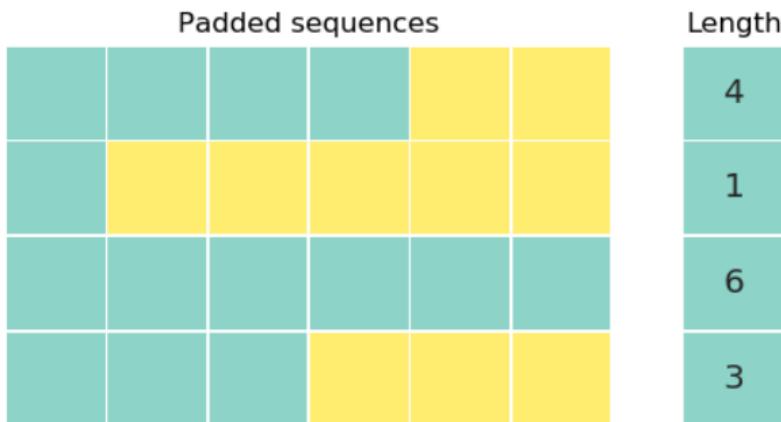
Position coding

Until now, Transformer does not take position into account !!



Masking attention computation

Think about **batches** :

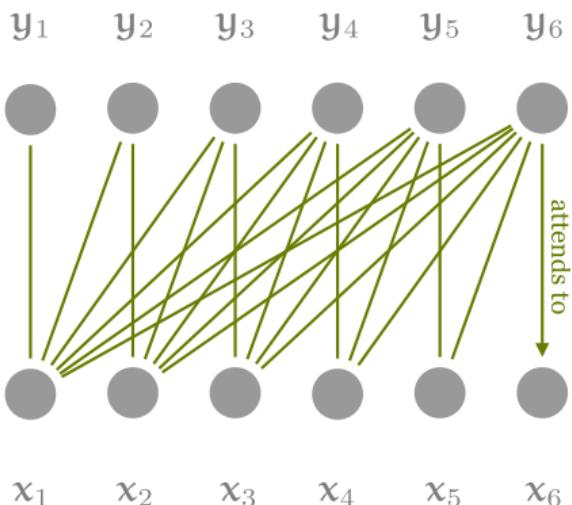
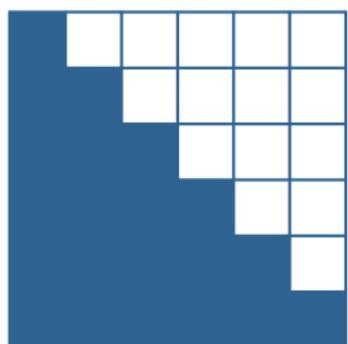
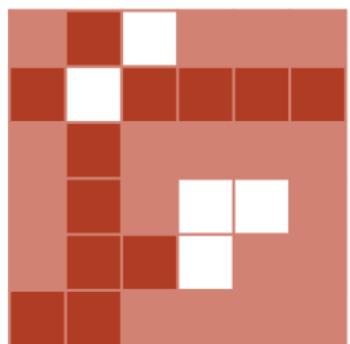


- Lengths are different
- Attention for PAD token = $-\infty \Rightarrow \alpha_{PAD} = 0$



Masking attention computation

Constraint on the attention to improve prediction



raw attention weights

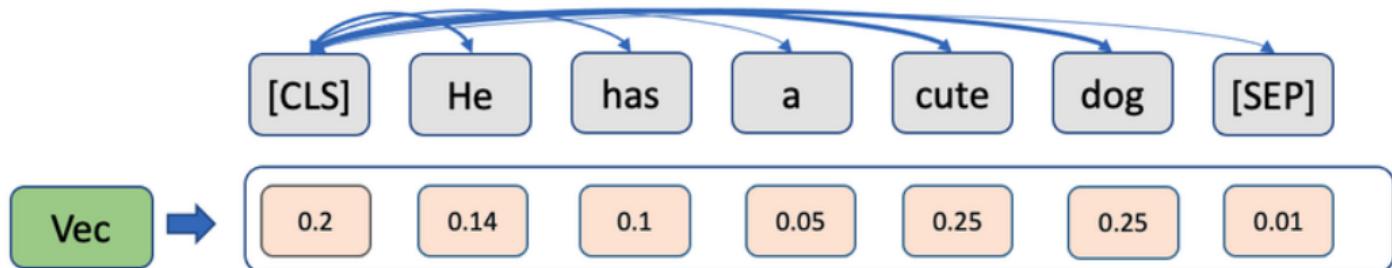
mask



Specific tokens

How to split texts or aggregate all informations? \Rightarrow Introducing specific tokens:

- CLS
- SEP



PRE-TRAINING / TRAINING TASKS

From W2V to MLM

- Masking
- Next token prediction (GPT)



Le docteur sorti du bloc [MASK] et...



But = prédire [MASK]

Next sentence

- Learn to aggregate... cheap manner

On chercher à savoir si une phrase en suit une autre

Input = [CLS] the man went to [MASK] store [SEP] he bought a gallon
[MASK] milk [SEP]

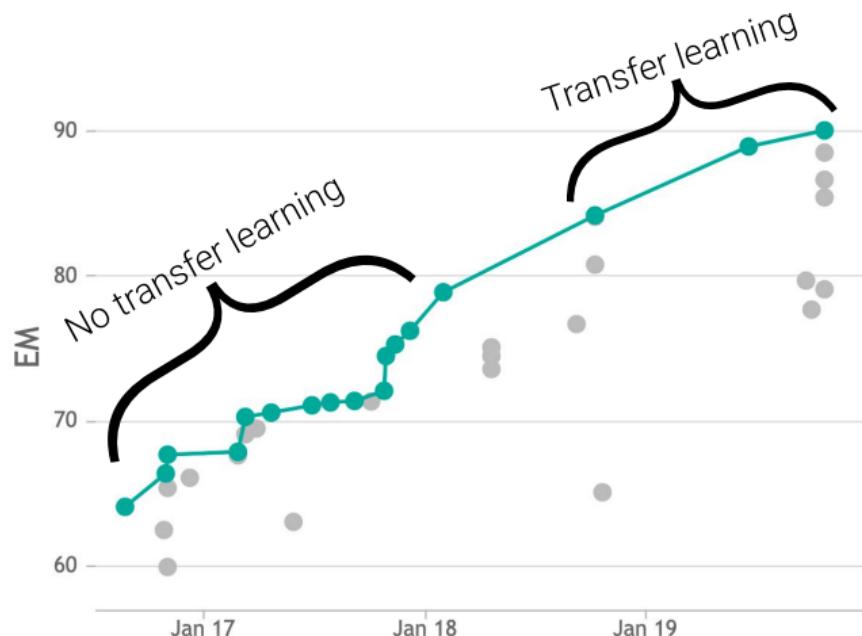
Label = IsNext

Input = [CLS] the man [MASK] to the store [SEP] penguin [MASK] are
flight ##less birds [SEP]

Label = NotNext

T5 architecture: generalizing gen. AI

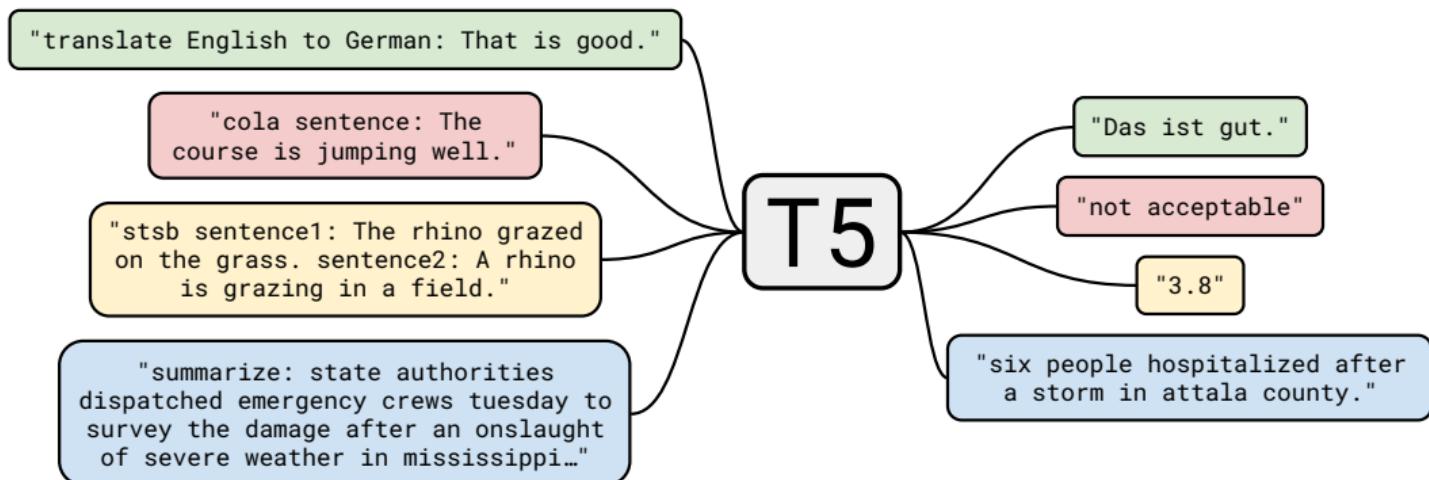
- Refining a generative model to answer different tasks ... T5 :)
- <https://colinraffel.com/talks/mila2020transfer.pdf>



Source: <https://paperswithcode.com/sota/question-answering-on-squad11-dev>

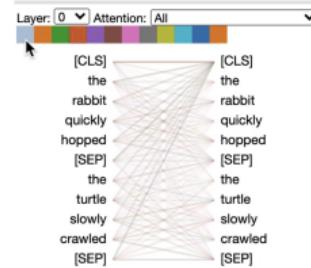
T5 architecture: generalizing gen. AI

- Refining a generative model to answer different tasks ... T5 :)
- <https://colinraffel.com/talks/mila2020transfer.pdf>



EXPLOITATION

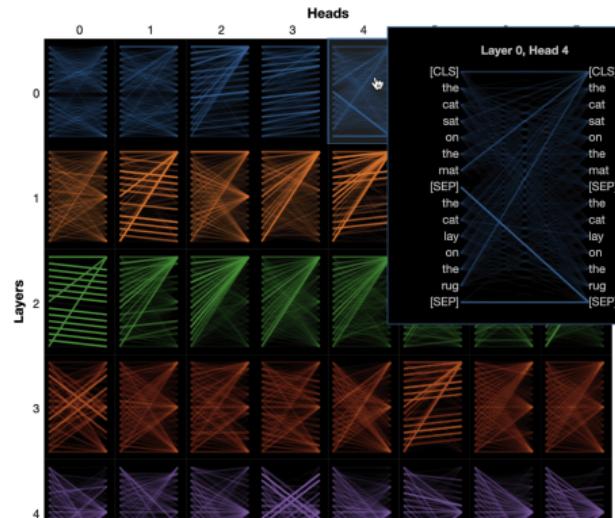
Layer visualization



Model View

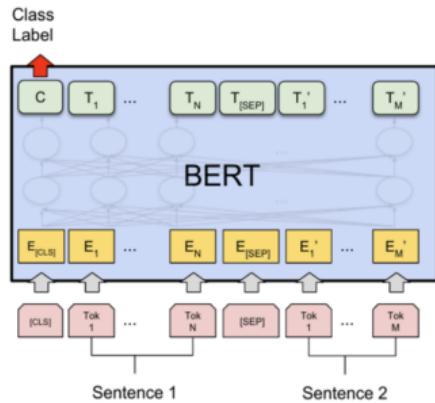
The *model view* shows a bird's-eye view of attention across all layers and heads.

Try out the model view in the [Interactive Colab Tutorial](#) (all visualizations pre-loaded).

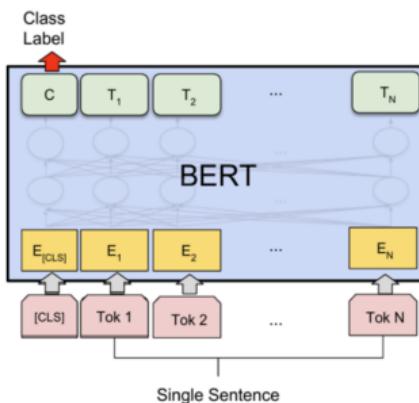


- BERT viz: <https://github.com/jessevig/bertviz>

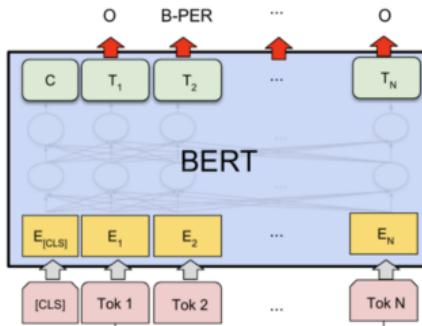
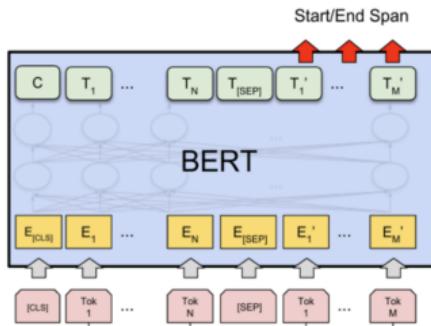
Exemple BERT



(a) Sentence Pair Classification Tasks:
MNLI, QQP, QNLI, STS-B, MRPC,
RTE, SWAG



(b) Single Sentence Classification Tasks:
SST-2, CoLA



Care of novelty: improvements are not guaranteed

- Time series